

Naitik Vora

(412) 512-4904 | naitik25vora@gmail.com | <https://www.linkedin.com/in/naitikvora/> | <https://github.com/NaitikVora>

EDUCATION

Carnegie Mellon University

Master of Information Systems Management

Courses: Data Structures & Algorithm, Artificial Intelligence, Database Systems, Distributed Systems and Cloud Computing

Pittsburgh, PA

Dec 2026

University of Mumbai

Bachelor of Engineering in Information Technology | GPA: 4.0/4.0

Mumbai, India

May 2023

SKILLS

Languages: Java, Python, Scala, Go, C++, C#, TypeScript, JavaScript, SQL, Bash

Databases: Oracle, MySQL, PostgreSQL, MongoDB, Elasticsearch, NoSQL Firebase DB, Google Real-time DB, Snowflake, Databricks, Redis

Frameworks and Tools: Spring Boot, Flask, GraphQL, ReactJS, Flutter, Docker, Kubernetes, Kafka, Cloud (AWS, GCP, Azure), LLMs

Others: Git, Jira, Jenkins, Confluence, Grafana, Agile, Linux, IoT, Splunk, Terraform, Prometheus, Operating systems, CoCo Agent SDK, Agentic AI

EXPERIENCE

Snowflake

Software Engineer Intern

Bellevue, WA

May 2026 - Aug 2026

- Owned the **Spark to Snowpark migration skills tooling**, leveraging predicate pushdown, column pruning, partition pruning, and cost-based query optimization to accelerate data retrieval and reduce distributed query execution time by **34%**
- Architected an **agent evaluation harness** with workload-specific evals and differential testing of legacy Spark and migrated Snowpark Connect outputs, achieving **96% validation accuracy**; and instrumented a Streamlit dashboard for regression diagnostics
- Engineered the LLM context-management layer: algorithmic LRU caching of prompt and translation artifacts plus context-overflow prevention via context slicing of large source files, eliminating token-limit failures and redundant LLM calls

PNC Bank (Carnegie Mellon University)

AI Engineer Fellowship

Pittsburgh, PA

Jan 2026 - Apr 2026

- Built **Brilliant Banker**, a full-stack AI banking assistant for small-business owners and PNC relationship managers, covering conversational guidance, LLM-generated client briefs, lead prioritization, and credit workflows
- Designed a **LangGraph** agent pipeline on the Anthropic Claude API with intent classification, tool routing, response generation, and **human-in-the-loop escalation gates** for high-stakes credit decisions

JPMorgan Chase

Software Engineer

Mumbai, India

Aug 2023 - Aug 2025

- Recognized with the JPMC Shining Star Award** for leading high-impact software delivery; strengthened cross-functional collaboration and improved team productivity by conducting **10+ technical knowledge-sharing sessions**
- Developed, optimized, and deployed 5+ high-performance Java Spring Boot microservices** supporting mission-critical banking operations; implemented RESTful APIs integrated with Oracle and MySQL, introduced connection-pool tuning, and improved query execution paths **raising transaction throughput by 20% and reducing production downtime by 15%** across global payment workflows
- Architected cloud-native backend systems** using AWS, Apache Kafka, and Gaia Kubernetes Platform (GKP), achieving **99.9% availability** during peak traffic and enabling auto-scaling to absorb **5x load spikes** without service degradation
- Led modernization of legacy financial systems** by migrating **3M+ records** to cloud infrastructure using Spring Batch with checksum validation and multi-threaded partitioning, ensuring **zero data loss**; decomposed a tightly coupled monolith into independently deployable microservices—**enhancing system scalability and reducing end-to-end API latency by 30%**
- Improved DevOps velocity and reliability** by enhancing CI/CD pipelines in Jenkins; automated build, test, and deployment workflows, integrated static code analysis & automated regression suites, and streamlined release orchestration—**cutting manual deployment errors by 70% and accelerating release/MVP cycles by 20%**
- Designed and implemented an AI-driven Retrieval-Augmented Generation (RAG) platform prototype**, integrating enterprise knowledge bases with LLMs to enable automated SQL query generation, debugging assistance, and incident-analysis workflows; **reduced investigation time by 30%** and demonstrated feasibility for firm-wide adoption of intelligent operational tooling

PROJECTS

[QuantQuest AI](#) | Python, pandas, NumPy, vectorbt, yfinance, Docker, ReactJS, Typescript, LLMs

Nov 2025

- Implemented AI-driven strategy evaluation pipelines**, integrating LLM inference, P&L computation engines, and risk-management logic to generate contextual coaching feedback and scenario-based analytics; added structured TypeScript models for reliability
- Designed an ML forecasting toolkit (scikit-learn, XGBoost) performing train/test splits, grid search tuning, k-fold cross-validation, and feature importance analysis; delivered improved price-movement prediction accuracy and automated reporting via Jupyter notebooks + Plotly dashboards

[Rural Health Tracker](#) | Python, Django, Flutter, TensorFlow, GCP

Aug 2023

- Engineered a cloud-deployed VGG19 CNN transfer-learning pipeline on Google Cloud** for medical image classification, achieving **96% accuracy** in detecting cataract and diabetic retinopathy
- Published award-winning IEEE paper** demonstrating technical impact: [“Early Diagnosis of Cataract and Diabetic Retinopathy using a Cloud-Based Deep Learning Model”](#) (ICACTA 2023), Mumbai, India, 2023, pp. 1-6, doi: 10.1109/ICACTA58201.2023.10392985